


What has been demonstrated

The September 10 Windows record supersedes the older pitch's preview status. These are retained results from the application work, not a new round of native tests performed for this portfolio.

Evidence	Recorded result
Windows preview	0.2.0 installed; source identity and fresh MCP handshake checked Windows / September 10, 2026
Automated checks	37 checks passed; Python compilation and JavaScript syntax passed Recorded source test run, Python 3.14
Native acceptance	11 extended checks plus two-design smoke and background/interactive tests Fusion 2704.1.53
Exact basic geometry	20 x 30 x 10 mm box: 6,000 mm3; radius 5 x 15 mm cylinder: 1,178.097 mm3 Native model measurements
Artifacts	Actual viewport PNGs and F3D/STEP/STL outputs; eight test artifacts retained privately Image/file identity checked
Managed Codex runs	GPT-6-Astra/low and native-picker GPT-5.6-Luna/low inspected the box; the latter captured a view and replied Recorded real CLI inference; not every model
Completed modeling workflow	12 solid components, 36 features and 36 sketches; native export/reopen and 12 STL exports followed by downstream QA One scoped workflow using advanced Python; no broad parity claim
Not yet verified	Mac native use, full Fusion restart/crash recovery, actual Claude/local inference and broad specialized CAD/CAM/simulation tools Open scope; no specialized trained model

AI-assisted workflows supported the creation of this project.

From a selected model to real geometry



Original native Windows screenshot. It shows the selected model, completed run, genuine box viewport and response. Incidental application labels are historical.

A native-picker run with GPT-5.6-Luna/low inspected a 20 x 30 x 10 mm box, obtained a viewport image and wrote a visible reply. A separate GPT-6-Astra/low runner test also inspected the box. This confirms those recorded paths, not every available model.

A subsequent modeling workflow used Formivect's opt-in advanced Python to create 12 native solids with 36 features and 36 sketches, validate native export/reopen and export 12 STLs. Downstream rendering, slicing and packaging were separate project tools, not additional built-in Formivect MCP features.

Control semantics. Pause holds future native work; Stop cancels queued jobs and ends managed inference at a control boundary. A current native operation can finish. Resume releases the queue; a new Start agent action begins another run. There is no automatic geometry rollback.

Remaining native checks. Mac execution, complete Fusion restart/crash recovery, live Claude/local inference and broader dedicated assembly/drawing/simulation/CAM operations remain open.

A measurable next workflow

A dimension-driven mounting-plate change order is the proposed benchmark candidate: revise hole spacing and thickness while preserving an interface datum and minimum edge distance, then inspect, export and verify a second manual edit.

Dedicated sketch constraints, hole-pattern tools, stable revision IDs and independent-native exports are explicit gaps. Advanced Python experiments do not substitute for a validated typed operation contract.

Gate	Evidence to advance
1 / Choose a useful workflow	Confirm a representative recurring task and expert-reviewable acceptance contract. The mounting-plate change order remains a candidate.
2 / Freeze independent evaluation	Hold out design families; define evaluator-owned geometry and editability oracles, errors and a full manual baseline.
3 / Make tools and recovery reliable	Build typed feature/constraint operations, stable targets, revision guards, checkpoints and interruption recovery.
4 / Establish model and data baselines	Compare licensed candidates; curate permissioned, executed trajectories with checked labels and provenance.
5 / Evaluate specialization	Train versioned candidates and compare unseen-task success, correction time, latency and resource cost; retain rollback.
6 / Validate independent operation	Integrate native engines, manual editing, persistent execution and owned inference into reviewed real-user workflows.
7 / Expand supported engineering	Add analyses, manufacturing, drawings and broader CAD with separate benchmarks, platform tests and production support.

Report completed/attempted tasks, requirement satisfaction, wrong-target edits, clarification quality, second-edit success, recovery, reviewer effort, latency and measured resource use. Hold out design families and keep the final evaluator inputs outside the agent workspace.

No customer adoption, revenue, completed independent platform, trained proprietary model or future delivery date is claimed.

AI-assisted workflows supported the creation of this project.